

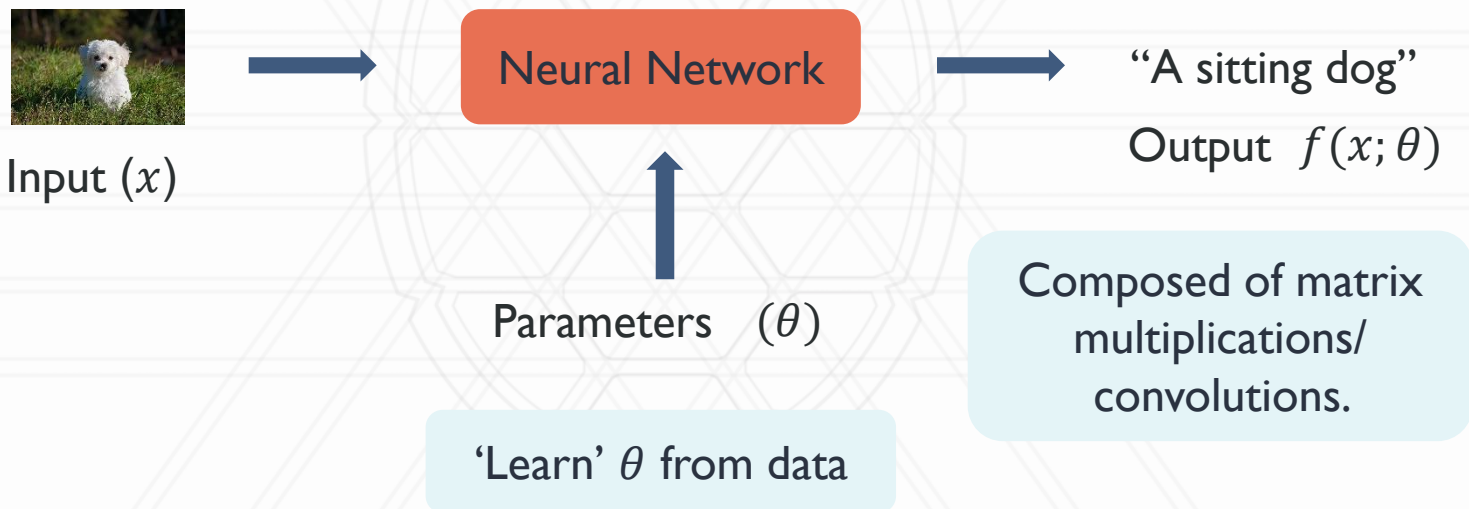
Tutorial. May 21, 2023 2-6  
pm

# Distributed Training of Deep Neural Networks

Abhinav Bhatele, Siddharth Singh  
Department of Computer Science

# Neural Networks


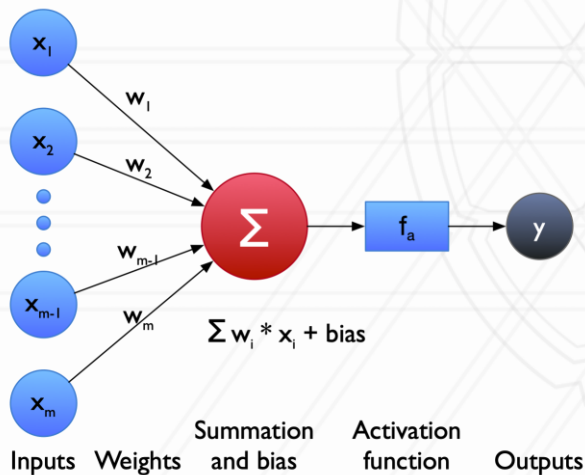
- Neural Networks (NNs): ‘Parameterized’ function approximators
- Can work with very high dimensional data.



# Deep neural networks

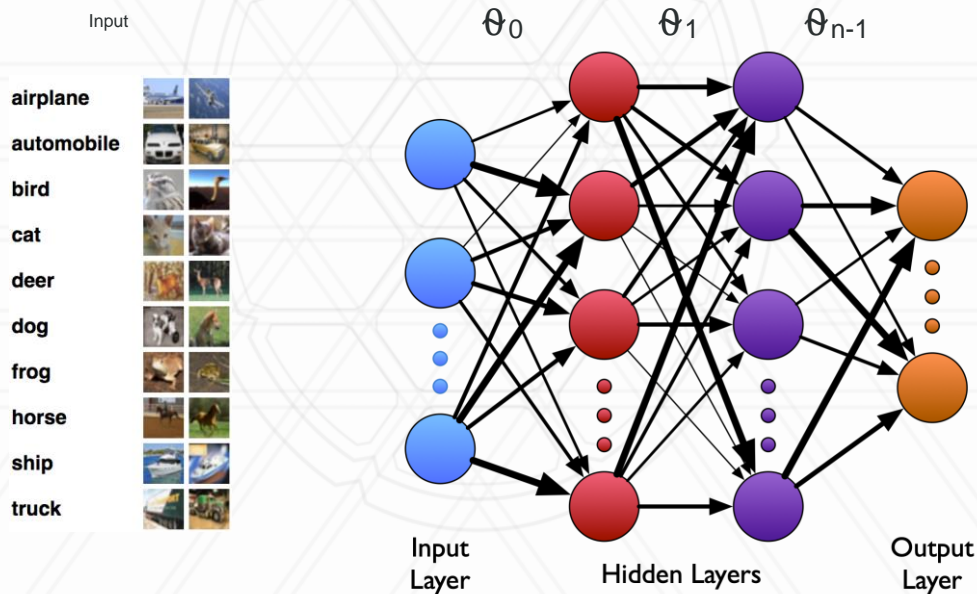
Neural networks can be used to model complex functions

Several layers that process input data



# Training a neural network

Problem: Find a set of weights/parameters that best fits the function we are trying to learn over a given training dataset



# Other terms and definitions

- Loss: a scalar proxy that when minimized leads to higher accuracy
- Learning/training: task of selecting weights that lead to an accurate function / minimizes the loss
- Gradient descent: process of updating the weights using gradients (derivatives) of the loss weighted by a learning rate
- Batch: Small subsets of the dataset processed independently
- Epoch: One pass over all the batches

# Stochastic Gradient Descent

Divide training data into batches

Repeat the following steps until loss,  $L$ , is minimized sufficiently:

- Read in one batch of training data
- Forward pass: Compute the activation,  $f(x; \theta)$ , and loss,  $L$ , on the batch
- Backward: Calculate gradients of the loss w.r.t. the parameters via backpropagation  $\frac{\partial L}{\partial \theta}$
- Optimizer step: Use gradients to update weights/parameters,  $\theta$ , such that loss is incrementally reduced

# Get the tutorial repository

---

- Clone the git repository as follows:

```
git clone https://github.com/hpcgroup/distrib-dl-tutorial.git
```

# PyTorch

---

- torch – a Python library for tensor computations with GPU support
- torch.nn – library for training deep neural networks
- We will start with looking at single GPU training using PyTorch



# Training task

- Image classification using MNIST data



# Using PyTorch

- Code location in the tutorial repo:  
session\_1\_basics/train.py

```
$ cd session_1_basics/  
$ sbatch --reservation=isc2023 run.sh
```

Parameter	
--num-layers	4
--hidden-size	2048
--image-size	64
--data-dir	<path-to-data>
--batch-size	32
--lr	0.001

# Mixed-precision Training

- GPUs have FP32, FP64 and tensor cores
- We can optimize performance by doing some operations in lower precision



# Mixed-precision Training

---

- Code location in the tutorial repo:  
session\_1\_basics/train\_mp.py

```
MIXED_PRECISION=true sbatch --reservation=isc2023 run.sh
```

# Activation Checkpointing

---

- Activations are outputs of individual layers
- To save memory, we checkpoint only inputs to each layer
  - Regenerate intermediate and output activations as needed in the backward pass
- Code location in the tutorial repo:  
`session_1_basics/train_mp.py`

```
CHECKPOINT_ACTIVATIONS=true sbatch --reservation=isc2023 run.sh
```



**Abhinav Bhatele and Siddharth Singh**

Department of Computer Science

[bhatele@umd.edu](mailto:bhatele@umd.edu), [ssingh37@umd.edu](mailto:ssingh37@umd.edu)

