

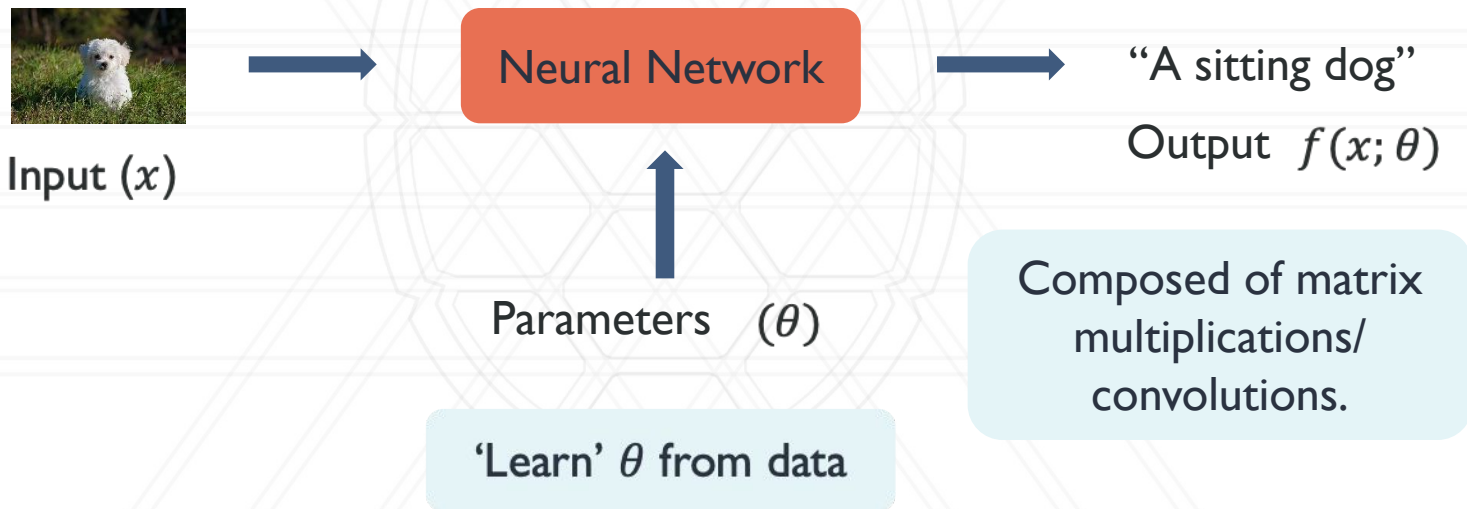
Tutorial. May 12, 2024 2-6 pm

Distributed Training of Deep Neural Networks

Abhinav Bhatele, Siddharth Singh, Daniel Nichols
Department of Computer Science

Neural Networks


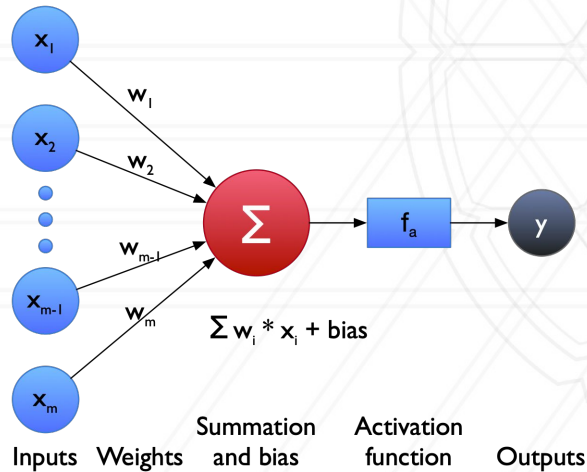
- Neural Networks (NNs): ‘Parameterized’ function approximators
- Can work with very high dimensional data.



Deep neural networks

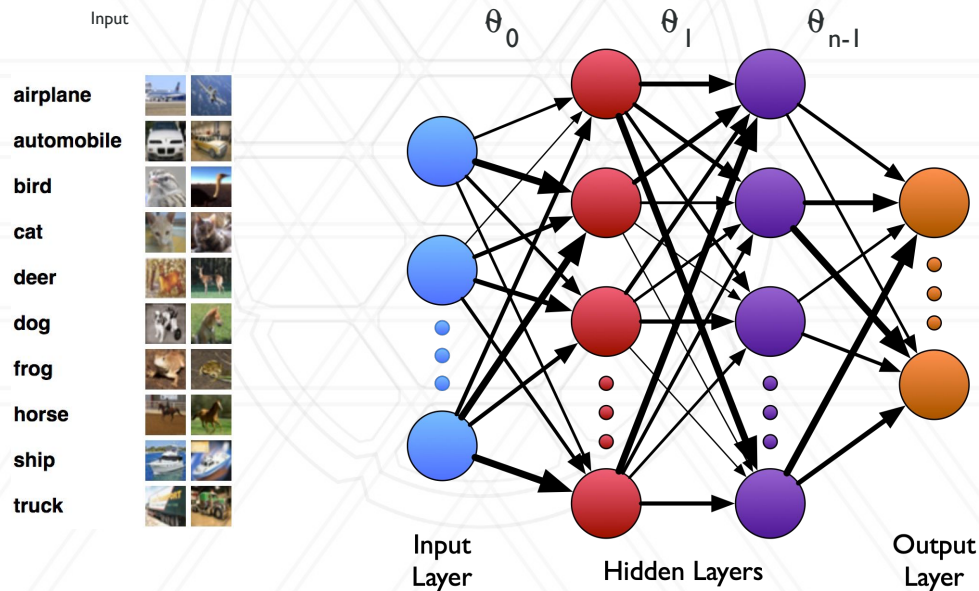
Neural networks can be used to model complex functions

Several layers that process input data



Training a neural network

Problem: Find a set of weights/parameters that best fits the function we are trying to learn over a given training dataset



Other terms and definitions

- Loss: a scalar proxy that when minimized leads to higher accuracy
- Learning/training: task of selecting weights that lead to an accurate function / minimizes the loss
- Gradient descent: process of updating the weights using gradients (derivatives) of the loss weighted by a learning rate
- Batch: Small subsets of the dataset processed independently
- Epoch: One pass over all the batches

Stochastic Gradient Descent

Divide training data into batches

Repeat the following steps until loss, L , is minimized sufficiently:

- Read in one batch of training data
- Forward pass: Compute the activation, $f(x; \theta)$, and loss, L , on the batch
- Backward: Calculate gradients of the loss w.r.t. the parameters via backpropagation $\frac{\partial L}{\partial \theta}$
- Optimizer step: Use gradients to update weights/parameters, θ , such that loss is incrementally reduced

Where are the matrix multiplies?

Embedding

Encoder

Encoder

...

Encoder

Classifier

Encoder

Linear

Self
attention

Linear

Attention block




Linear

ReLU

Linear

Multi-layer perceptron

Linear



Get the tutorial repository

- Clone the git repository as follows:

```
git clone https://github.com/axonn-ai/distrib-dl-tutorial.git
```


PyTorch

- torch – a Python library for tensor computations with GPU support
- torch.nn – library for training deep neural networks
- We will start with looking at single GPU training using PyTorch

Training task

- Image classification using MNIST data



Using PyTorch

- Code location in the tutorial repo: `session_1_basics/train.py`

```
$ cd session_1_basics/  
$ sbatch --reservation=isc2024 run.sh
```

Parameter	
--num-layers	4
--hidden-size	2048
--image-size	64
--data-dir	<path-to-data>
--batch-size	32
--lr	0.001

Mixed-precision Training

- GPUs have FP32, FP64 and tensor cores
- We can optimize performance by doing some operations in lower precision



Mixed-precision Training

- Code location in the tutorial repo:
`session_1_basics/train_mp.py`

```
MIXED_PRECISION=true sbatch --reservation=isc2024 run.sh
```

Activation Checkpointing

- Activations are outputs of individual layers
- To save memory, we checkpoint only inputs to each layer
 - Regenerate intermediate and output activations as needed in the backward pass
- Code location in the tutorial repo:
`session_1_basics/train_mp.py`

```
CHECKPOINT_ACTIVATIONS=true sbatch --reservation=isc2024 run.sh
```



Abhinav Bhatele and Siddharth Singh

Department of Computer Science

bhatele@umd.edu, ssingh37@umd.edu

