

# Distributed Deep Learning on GPU-based Clusters

Abhinav Bhatele, Siddharth Singh, Prajwal Singhanian  
Department of Computer Science


# Limitations of data parallelism

---

- DDP – Supports models of limited size
- FSDP – Large communication overheads, specially for small batch size tasks like IFT

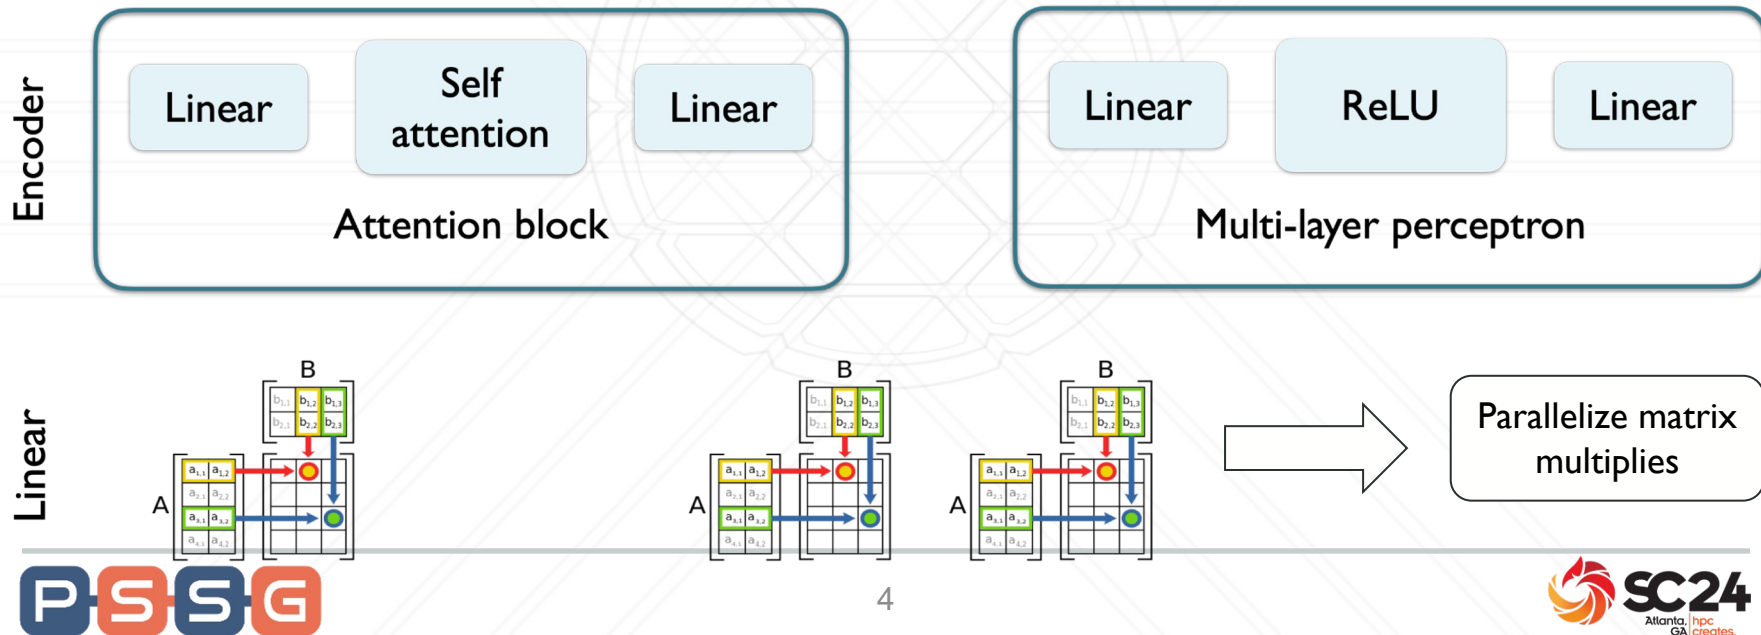
# Tensor parallelism

- Divide parameters and compute of every layer of a neural network on multiple GPUs.




# Tensor parallelism

- Divide parameters and compute of every layer of a neural network on multiple GPUs.



# AxoNN's 3D Tensor Parallelism



Parallelizing a matrix multiplication ( $I \cdot W = O$ ) using AxoNN on 8 GPUs

# Creating an AxoNN Lightning Strategy

```
from axonn.lightning import AxonnStrategy
```

```
pl_strategy = AxonnStrategy(
```

```
    G_intra_x=..
```

```
    G_intra_y=..
```

```
    G_intra_z=..
```

```
    overlap_communication=True,
```

```
)
```

3D tensor parallel grid  
dimensions

# Running the code (Tensor/Intra-Layer)

---

- Code – `train.py`

```
CONFIG_FILE=configs/axonn.json  
sbatch --ntasks-per-node=4  train.sh
```

# Let's try different AxoNN configurations

---

- In config/axonn.json - tp\_dimensions is [2, 2, 1].
- Now change it to [4, 1, 1] and rerun your code.
- Does it become faster?



# How to set the parallel configuration?

---

- Small batch-sizes - Use x and/or y
  - Example - finetuning and inference
- Large batch sizes - Use z (+data parallelism)
  - Example - pretraining

# Alternative way to use AxoNN

---

```
from axonn.intra_layer import auto_parallelize
with auto_parallelize():
    net = FC_Net(args.num_layers, args.image_size**2, args.hidden_size, 10).cuda()
```

Zero code changes required in your model definition!

`fabric.init_module()` calls AxoNN's auto-parallelize function